

NoSQL databases

Dmitry A. Duev

NoSQL databases

NoSQL: not only SQL / non SQL

- What drove the development of NoSQL?
 - Pros and cons of traditional SQL systems
- What are the different flavors of NoSQL?
 - Wide column-based
 - Key-value
 - Document-based
 - Graph-based

Pros and cons of traditional RDBMS

+

- Have been around for quite a while
- SQL: mature and powerful
- Transactions + ACID-compliance built-in
- Data normalization + joins
- Schema: good when data can be represented in appropriate way
- Many open-source solutions

-

- Vertical scale-out | scaling issues
- Joins + transactions across multiple dbs quickly become costly for complicated objects and big data -> performance and availability are affected
- [Object-relational impedance mismatch](#)
- Schema: bad for unstructured/evolving data

NoSQL: motivation and promise

The design and development of NoSQL databases has been largely driven by the RDBMS cons from the previous slide.

- Deliver performance for big, potentially unstructured or evolving data that may come in in real time
 - Simple design
 - [Horizontal scaling](#)
- NoSQL ~= not only SQL, i.e. some systems support SQL-like query languages

NoSQL: fulfilling the promise

Always comes with a cost!

- Most such systems lack ACID transactions and offer BASE (Basic Availability, Soft state and Eventual consistency) instead
 - See [ACID vs BASE](#)
 - Some systems exhibit potential lost writes and other forms of data loss
 - A notable exception is MongoDB
- No schema means data integrity might become an issue
 - Some systems do allow defining schemas and perform validation/enforcement (e.g. MongoDB)
- Many query languages vs single SQL (albeit with different flavors)

NoSQL: fulfilling the promise

In practice, you almost always still need to deal with relational data! There are three main techniques to do that:

- Nesting/embedding data
 - Store all data needed for a specific task in one place (e.g. in a single document)
- Linking + Multiple queries
 - Store a foreign key and fetch data in multiple queries. Since single queries in NoSQL are often more performant than in SQL, may be ok
- Caching and replication
 - Instead of storing a foreign key, store the actual values

Data schema modelling must be done very differently from RDBMS. A simple translation would often not work.

Main types of NoSQL databases

- Key-value store
 - Uses maps/dictionaries/associated lists/hash tables with [corresponding operation complexities](#)
 - Examples: [Redis](#), [ArangoDB](#), [ZooKeeper](#), [Couchbase](#), [Cassandra](#), [Amazon DynamoDB](#)
- Wide column store
 - Essentially, a two-dimensional semi-structured key-value store
 - Examples: [Cassandra](#), [HBase](#)
- Document store
 - Semi-structured; data are encapsulated in some standard form (XML, JSON, BSON) in “documents” with unique keys/identifiers
 - Often uses B-trees with [corresponding operation complexities](#)
 - Examples: [Couchbase](#), [MongoDB](#), [Amazon DynamoDB](#)
- Graph store
 - Uses graphs to represent data + relationship between them (the latter can be queried, too)
 - Examples: [Neo4J](#), [ArangoDB](#)

Performance

Data model	Performance	Scalability	Flexibility	Complexity	Functionality
Key-value store	high	high	high	none	variable
Column-oriented store	high	high	moderate	low	minimal
Document-oriented store	high	variable (high)	high	low	variable
Graph database	variable	variable	high	high	graph theory
Relational database	variable	variable	low	moderate	relational algebra